

Performance Challenges of DeepLearning

- Compute intensive process



Solutions

- Move from CPUs to GPUs
- Parallelism
- Shard the dataset
- Train on Multiple Cores
- Train on Multiple Machines



Deep Learning Review

- Simple Neural Network
- Classify Flowers by measurements



The Data

5.1,3.5,1.4,0.2,0
4.9,3.0,1.4,0.2,0
4.7,3.2,1.3,0.2,0
4.6,3.1,1.5,0.2,0



The Code

- <https://github.com/deeplearning4j/dl4j-examples/blob/master/dl4j-examples/src/main/java/org/deeplearning4j/examples/dataexamples/CSVExample.java>



Code Overview

- Read the Data
- Build a model
- Train the model

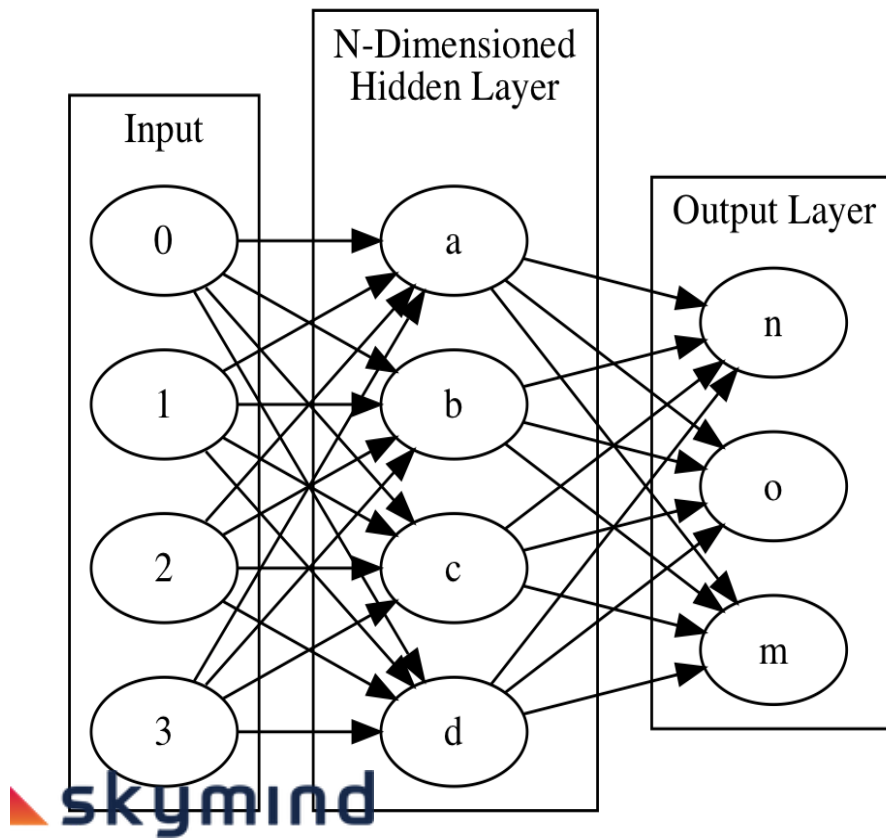


Performance Data Iris Model

- read 150 lines 3KB
- Train the model 39 parameters



Picture of Model Graph



VGG16

- Read 14,197,122 images
- Augment with image transformations
- Manage 13,8357,544 parameters
- 21 Layers
- Train for 72 epochs



Picture of VGG



Options for Increasing Performance



Options for executing code

- CPU
- GPU
- Multi GPU
- Spark
- SPark + GPU



Architecture Details

- libnd4j
 - The C++ engine that powers ND4J
- ND4J
 - n-dimensional arrays for java
- Datavec
 - ETL
- DeepLEarning4j
 - Build, train and deploy models



ND4J

- Where the math happens
- ND4J-backend
- Determines how execution happens



Execute on CPU

```
<Your pom.xml file>  
<properties>  
<nd4j.backend>nd4j-native-platform</nd4j.backend>
```



Execute on GPU

```
<Your pom.xml>  
<properties>  
<nd4j.backend>nd4j-cuda-8.0-platform</nd4j.backend>
```



Troubleshooting

- Failure to use multiple GPU's

```
//Add this to your main  
CudaEnvironment.getInstance().getConfiguration()  
    .allowMultiGPU(true);
```



Multi-GPU Data Parallelism

- ParallelWrapper
 - Model Config remains the same
 - ND4J-Backend set to GPU

```
ParallelWrapper wrapper = new ParallelWrapper.Builder(YourExistingModel)  
    .prefetchBuffer(24) .workers(4) .averagingFrequency(1) .reportScoreAfterAveraging(true)  
    .useLegacyAveraging(false) .build();
```



ParallelWrapper Internals

- Model will be duplicated
- Each worker trains model
- Models averaged at averagingFrequency(X)
- Training Continues



Distributed Training in Spark

- Documentation
 - <https://deeplearning4j.org/spark>
- Examples
 - <https://github.com/deeplearning4j/dl4j-examples/tree/master/dl4j-spark-examples>



Training on Spark

- SparkDl4jMultiLayer vs MultiLayerNetwork
- SparkComputationGraph vs ComputationGraph



Non-Distributed Training overview

- Process minibatch
 - Calculate Loss Function
 - Calculate direction of less error
 - Update weights in that direction



Non-Distributed Training Challenges

- Number of parameters may be large
- CPU or preferably GPU intensive
- Update large multi-dimensional matrix of numeric values



Data Parallelism vs Model Parallelism

- Data Parallelism

Different machines have a complete copy of the model; each machine simply gets a different portion of the data, and results from each are combined.

- Model Parallelism

different machines in the distributed system are responsible for the computations in different parts of a single network - for example, each layer in the neural network may be assigned to a different machine.



Data Parallelism

- Our current approach



Data Parallel Theoretical options

- Parameter averaging vs. update (gradient)-based approaches
- Synchronous vs. asynchronous methods
- Centralized vs. distributed synchronization



DL4J Current approach

- Synchronous Parameter Averaging



Parameter Averaging Overview


1. Initialize the network parameters randomly based on the model configuration
2. Distribute a copy of the current parameters to each worker
3. Train each worker on a subset of the data
4. Set the global parameters to the average the parameters from each worker
5. While there is more data to process, go to step 2



Distributed Training Implementation

- Workers process minibatch
- Calculate Loss Function
- Calculate Gradient update after each minibatch
- Submit to Parameter server
- Parameter Server averages weights from workers
- Ships averaged weights to workers





Parameter Averaging Considerations

- Distributed Systems 101
- Synchronization is a challenge to performance
- Sharing is a challenge to performance



Tuning Parameter Averaging

- Setting Averaging Period
- Increase and workers may diverge
- Decrease and cost of Synching increases
- Advice
 - Greater than 1
 - Less than 20
 - Between 5-10



Batch Prefetch

Spark workers are capable of asynchronously prefetching a number of minibatches (DataSet objects), to avoid waiting for the data to be loaded.

- Values
 - 0 disables prefetching.
 - 2 is often a sensible default.
 - Too Large and Memory is wasted with no additional benefit



Updater State

- momentum, RMSProp and AdaGrad have internal state
- saveUpdater == true
 - updater state (at each worker) will be averaged and returned to the master along with the parameters
 - Extra time and Network Traffic, may improve performance



More Options

- See <https://deeplearning4j.org/spark>

DL4J Spark Examples

- <https://github.com/deeplearning4j/dl4j-examples/tree/master/dl4j-spark-examples>



Minimal Example

```
JavaSparkContent sc = ...;
JavaRDD<DataSet> trainingData = ...;
MultiLayerConfiguration networkConfig = ...;

//Create the TrainingMaster instance
int examplesPerDataSetObject = 1;
TrainingMaster trainingMaster = new ParameterAveragingTrainingMaster.Builder(examplesPerDataSetObject)
    .(other configuration options)
    .build();

//Create the SparkDl4jMultiLayer instance
SparkDl4jMultiLayer sparkNetwork = new SparkDl4jMultiLayer(sc, networkConfig, trainingData);

//Fit the network using the training data:
sparkNetwork.fit(trainingData);
```



How to Participate and Contribute

- Chat with us on Gitter
 - <https://gitter.im/deeplearning4j/deeplearning4j>
- Contribute
 - <https://github.com/deeplearning4j/>



Interesting Challenges

- GPU aware Yarn
 - <https://issues.apache.org/jira/browse/YARN-5517>
- Parallelism in DeepLearning
 - https://static.googleusercontent.com/media/research.google.com/en//archive/large_deep_netwc




Early Stopping

- Allows Training to stop before overfitting occurs



Early Stopping Image



ParallelWrapper and Early Stopping

- What is Early Stopping?
 - Saves Model that performs best on Test Data
 - Saves Model before training overfits training data



Training on Single GPU or CPU

- Use EarlyStoppingTrainer



Training on Multi-GPUs

- Use `EarlyStoppingParallelTrainer`



Training on Spark

- Use SparkEarlyStoppingTrainer



Examples



Main Examples Repo

<https://github.com/deeplearning4j/dl4j-examples>



Quick Start Guide

- <https://deeplearning4j.org/quickstart>



Spark Examples

- <https://github.com/deeplearning4j/dl4j-examples/tree/master/dl4j-spark-examples>
- Includes
 - EMR
 - Distributed Spark MLP example
 - LSTM on Spark Example
 - Statistics Gathering Example



GPU Examples

- <https://github.com/deeplearning4j/dl4j-examples/tree/master/dl4j-cuda-specific-examples>



Material included for home study

The following materials are included for your use, not intended to be covered during presentation



DataVec: Spark Transform Process



DataVec overview

- Goal of DataVec
- Data => INDArray



DataVec Spark Based Operations

- Joins
- Column Transformations
- Data Analysis



Spark Analyze

- Tool to gather statistics across a data set



Spark Analyze Example

- The Data: classic dataset of Iris flower characteristics

```
5.1,3.5,1.4,0.2,0  
4.9,3.0,1.4,0.2,0  
4.7,3.2,1.3,0.2,0  
4.6,3.1,1.5,0.2,0  
5.0,3.6,1.4,0.2,0  
5.4,3.9,1.7,0.4,0
```



Spark Analyze Example Continued...

- Define a Schema

```
Schema schema = new Schema.Builder()  
    .addColumnDouble("Sepal length", "Sepal width",  
        "Petal length", "Petal width")  
    .addColumnInteger("Species")  
    .build();
```



Spark Analyze Example Continued...

- Create Spark Conf

```
SparkConf conf = new SparkConf();  
conf.setMaster("local[*]");  
conf.setAppName("DataVec Example");  
JavaSparkContext sc = new JavaSparkContext(conf);
```



Spark Analyze Example Continued...

- Read Data

```
String directory = new ClassPathResource("IrisData/iris.txt")
    .getFile().getParent();
//Normally just define your directory
//like "file:/" or "hdfs:/"
JavaRDD<String> stringData = sc.textFile(directory);
```



Spark Analyze Example Continued...

- Parse Data

```
//We first need to parse this comma-delimited (CSV) format;  
//we can do this using CSVRecordReader:  
RecordReader rr = new CSVRecordReader();  
JavaRDD<List<Writable>> parsedInputData =  
stringData.map(new StringToWritablesFunction(rr));
```



Spark Analyze Example Continued...

- Analyze Data

```
int maxHistogramBuckets = 10;
DataAnalysis dataAnalysis =
    AnalyzeSpark.analyze(schema,
        parsedInputData, maxHistogramBuckets);

System.out.println(dataAnalysis);
```



Spark Analyze Output

- min,max,mean,StDev,Variance,num0,
- numNegative,numPositive,numMin,numMax,Total
- See example IrisAnalysis.java



Spark Analyze options

- Per Column Analysis

```
//We can get statistics on a per-column basis:  
DoubleAnalysis da = (DoubleAnalysis)dataAnalysis.getColumnAnalysis("Sepal"  
double minValue = da.getMin();  
double maxValue = da.getMax();  
double mean = da.getMean();
```



Spark Analyze HTML output

- To Generate HTML Output

```
HtmlAnalysis.createHtmlAnalysisFile(dataAnalysis,  
new File("DataVecIrisAnalysis.html"));  
  
//To write to HDFS instead:  
//String htmlAnalysisFileContents =  
//HtmlAnalysis.createHtmlAnalysisString(dataAnalysis);  
//SparkUtils.writeStringToFile("hdfs://your/hdfs/path/here",  
//htmlAnalysisFileContents,sc);
```



DataVec Spark Transform Process

- Joining
- Transformation
- Schema alteration



Video

- <https://www.youtube.com/watch?v=MLEMw2NxjxE>



Code Examples

- <https://github.com/deeplearning4j/dl4j-examples/tree/master/datavec-examples>
- http://github.com/SkymindIO/screencasts/tree/master/datavec_spark_transform



Basic Example CSV Data

- Storm Reports Data

```
161006-1655,UNK,2 SE BARTLETT,LABETTE,KS,37.03,-95.19,  
TRAINED SPOTTER REPORTS TORNADO ON THE GROUND. (ICT),TOR
```

```
//Fields are
```

```
//datetime,severity,location,county,state,lat,lon,comment,type
```



Define Schema as read

```
Schema inputDataSchema = new Schema.Builder()  
    .addColumnString("datetime", "severity", "location", "county", "state")  
    .addColumnDouble("lat", "lon")  
    .addColumnString("comment")  
    .addColumnCategorical("type", "TOR", "WIND", "HAIL")  
    .build();
```



Define Transform process

```
TransformProcess tp = new TransformProcess.Builder(inputDataSchema)
    .removeColumns
    ("datetime","severity","location","county","state","comment")
    .categoricalToInteger("type")
    .build();
// keep lat/lon convert type to 0,1,2
```



Create Spark Conf

```
SparkConf sparkConf = new SparkConf();  
sparkConf.setMaster("local[*]");  
sparkConf.setAppName("Storm Reports Record Reader Transform");  
JavaSparkContext sc = new JavaSparkContext(sparkConf);
```



read the data file

```
JavaRDD<String> lines = sc.textFile(inputPath);
```



Convert to Writable

```
JavaRDD<List<Writable>> stormReports =  
lines.map(new StringToWritablesFunction(new CSVRecordReader()));
```



Run our transform process

```
JavaRDD<List<Writable>> processed =  
SparkTransformExecutor.execute(stormReports, tp);
```



Convert Writable Back to String for Export

```
JavaRDD<String> toSave=  
processed.map(new WritablesToStringFunction(", "));
```



Write to File

```
toSave.saveAsTextFile(outputPath);
```

